


MA Project: MADDNESS Accelerator

Jannis Schönleber, Matteo Perotti, Lukas Cavigelli, Renzo Andri



Motivation

- **Matrix multiplications are key:**
 - Machine learning
 - Numerical solvers
- **New method for approximate matrix multiplications (“MADDness”): ICML’21, [arXiv](#)**

Multiplying Matrices Without Multiplying

Davis Blalock^{1 2} John Gutttag²

¹MosaicML, San Francisco, CA, USA ²MIT CSAIL

- code available: python ([here](#), [here](#)) and C ([here](#), with AVX2)
- **Convenient for many applications:**
 - not just strong quantization (fewer issues): continuous, more stable/less noisy
 - applicable to anything approximate or with iterative refinement
 - more fine-grained & flexible than quantization

MADDness: Method

What does MADDness solve?

- $A \in \mathbb{R}^{N \times D}, B \in \mathbb{R}^{D \times M}, N \gg D \geq M$
- Given compute time budget τ , find $g(\cdot), h(\cdot), f(\cdot), \alpha, \beta$, such that

$$\|\alpha f(g(A), h(B)) + \beta - AB\|_F < \varepsilon(\tau) \|AB\|_F$$

MADDNESS:

- makes use of B matrix being fixed (“weights”)
- input for inference:
 - (1) matrix A ,
 - (2) tensor of look-up tables T computed from B
- $(AB)_{n,m} \approx \alpha f(g(A), h(B))_{n,m} + \beta = \alpha \sum_{c=1}^C T_{m,c,k} + \beta$


with $k = g^{(c)}(a_n)$

ETH zürich

not a full DNN, just last layer(s)!

use VGG-like DNN for FE,
then learn 512→1000
linear classifier,
then approximate this

based on 1-thread CPU
implementation



Algorithm 1 MADDNESSHASH

```

1: Input: vector  $x$ , split indices  $j^1, \dots, j^4$ , split thresholds  $v^1, \dots, v^4$ 
2:  $i \leftarrow 1$  // node index within level of tree
3: for  $t \leftarrow 1$  to 4 do
4:    $v \leftarrow v_i^t$  // lookup split threshold for node  $i$  at level  $t$ 
5:    $b \leftarrow x_{j^t} \geq v ? 1 : 0$  // above split threshold?
6:    $i \leftarrow 2i - 1 + b$  // assign to left or right child
7: end for
8: return  $i$ 
    
```

Compute Workload

- **Formula:**

$$(AB)_{n,m} \approx \alpha f(g(A), h(B))_{n,m} + \beta = \alpha \sum_{c=1}^C T_{m,c,k} + \beta \quad \text{with } k = g^{(c)}(a_n)$$

with $A \in \mathbb{R}^{N \times D}$, $B \in \mathbb{R}^{D \times M}$, C the #codebooks, K the #learned prototypes

- **Computation & Lookups**

- $g(A)$: $\Theta(NC)$
- $h(B)$: $\Theta(MKCD)$, typically offline
- $f(\cdot, \cdot)$: $\Theta(NCM)$, with C table lookups for each output value (M cols, N rows)
- overall: $\Theta(MC(KD + N)) \xrightarrow{K=16, N \gg D} \Theta(NCM)$

- **Storage** (based on the example in the paper)

- VGG-style, CIFAR-100, classifier: $D=512$, $K=16$, $M=100$, $C \in \{16, 32, 64\}$
- $T \in \mathbb{R}^{M \times C \times K}$, thus in the order of 100kB ($C=32$, float16)
- reference: the “normal” linear layer would also use $D \cdot M$ weights, i.e. 100kB

Project Plan (MA: 26 weeks)

- **Apply Maddness to 1x1 conv layers in ResNet-50 [3-4 weeks]**
 - get familiar with method
 - define/specify baseline quantized NN
 - replace some conv1x1 layers and find good parameters
 - collect statistics on access patterns, parameters (tables sizes, hash params, locality)
- **System & device-level architecture exploration [4-6 weeks]**
 - analyze perf/power/area trade-offs and flexibility for fundamentally different architecture, tune the memory hierarchy & interconnect, and *consider load balancing* (possibly with software interaction):
 - fully-specialized systolic arrays (also chiplets, multi-chip scale-out)
 - a PULP-style (MemPool-style?) architecture with ISA extension and/or per-cluster accelerator(s)
 - answer fundamental implementation questions: Flexible or fixed size matmul? Store values or hashes for FMs?
 - identify key building blocks & refine their PPA estimates
- **Implement (RTL & backend), integrate, verify, evaluate the most promising architecture [6-7 weeks]**
- **Develop basic software to feed & control the circuit [4-5 weeks]**
- **Map a full ResNet-50 or a recent BERT or MobileNetV3 model, think about other applications [4-5 weeks]**
- **Report, presentation, final experiments, code clean-up [3 weeks]**

ETH zürich



HUAWEI

ZURICH **RESEARCH** CENTER

Weekly Update 0

- What happened
 - Beat COVID 😊 started working on Sunday
 - Git setup + remote setup (e.g. Badile, Picos, condor...)
 - Python code cleanup 17k -> 1.3k (right now) -> 800 LOC possible

Weekly Update 0

- **Python Cleanup**

- Merged all madness code
- Pylint + Pytest + Github CI integration for easy and fast refactoring

- **GOAL:**

- Close to our + paper mathematical formulation (bolt code far away)
- Have a good structure for exploration + collect statistics on access patterns, parameters (tables sizes, hash params, locality)

- **Mathematical formulation:**

$$(AB)_{n,m} \approx \alpha f(g(A), h(B))_{n,m} + \beta = \alpha \sum_{c=1}^C T_{m,c,k} + \beta \text{ with } k = g^{(c)}(a_n)$$

with $A \in \mathbb{R}^{N \times D}$, $B \in \mathbb{R}^{D \times M}$, C the #codebooks, K the #learned prototypes

Weekly Update 0

```
XtX = XtX.astype(np.float64)
XtY = XtY.astype(np.float64)

# preconditioning to avoid numerical issues (seemingly unnecessary, but
# might as well do it)
# scale = 1. / np.std(XtX)
if precondition:
    # # pretend cols of X were scaled differently
    # xscales = np.linalg.norm(XtX, axis=0) + 1e-20
    # mulby = (1. / xscales)
    # XtX *= mulby * mulby
    # XtY *= mulby.reshape(-1, 1)

    # yscale = np.linalg.norm(XtY, axis=1) + 1e-20
    # yscale = np.linalg.norm(XtY, axis=0) + 1e-20
    # yscale = yscale.reshape(-1, 1)

    # xscales = np.mean(np.linalg.norm(XtX, axis=0))
    # xscales = 7
    # xscales = 1

    # XtY *= (1. / yscale)
    # XtY *= (1. / yscale.reshape(-1, 1))

    # scale = 1. / len(X_enc)
    scale = 1.0 / np.linalg.norm(XtX, axis=0).max()
    XtX = XtX * scale
    XtY = XtY * scale

# W = np.linalg.solve(XtX, XtY)
W, _, _, _ = np.linalg.lstsq(XtX, XtY, rcond=None) # doesn't fix it

# W, _, _, _ = np.linalg.lstsq(X_bin, Y, rcond=None)

# import torch
# import torch.nn.functional as F
# import torch.optim as optim

# def _to_np(A):
#     return A.cpu().detach().numpy()
```

```
def optimal_split_val(
    self, X, dim, possible_vals=None, X_orig=None, return_possible_vals_losses=False
):
    if self.N < 2 or self.point_ids is None:
        if return_possible_vals_losses:
            return 0, 0, np.zeros(len(possible_vals), dtype=X.dtype)
        return 0, 0
    my_idxs = np.asarray(self.point_ids)
    if X_orig is not None:
        X_orig = X_orig[my_idxs]
    return optimal_split_val(
        X[my_idxs],
        dim,
        possible_vals=possible_vals,
        X_orig=X_orig,
        return_possible_vals_losses=return_possible_vals_losses,
    )
```

```
# I honestly don't know why this is the formula, but wow
# does it work well
bias = self.number_of_codebooks / 4 * np.log2(self.upcast_every)
dists -= int(bias)
```

Weekly Update 0

- **Next steps**
 - Finish python refactoring
 - C++ code refactoring
 - Build ResNet-50 Model with Tensorflow or PyTorch and integrate Maddness
- **Ideas & Discussion**
 - Current project/git repo name (uninspired) is halutmatmul
 - Is it possible to find a general hash/encoding (splits, thresholds) function with $C=..$ $K=..$ for AI application with batch normalized or normalized data? We would not need to learn those.

Research Questions

1. Can this be used for DNN inference (ResNet, BERT [more sensitive, maybe focus on matmuls with Q, K, V])? What are typical parameters? What are typical access patterns?
2. What is the most efficient (energy, throughput/area) hardware architecture we can think of for this? What is a suitable memory hierarchy, memory types?
3. What is the more general design/trade-off space? Are there hard cliffs? Can we make it flexible/scalable/more efficient/... as a system?
4. What is a minimal change that gets most of the benefits (e.g., ISA extension)?
5. How efficient is it with an actual implementation?
6. Are there more good applications/use-cases? Also training? Higher accuracy with quantization-aware training?